

# Information Retrieval

Weike Pan

The slides are **adapted from those provided by Prof. Hinrich Schütze** at University of Munich (<http://www.cis.lmu.de/~hs/teach/14s/ir/>).

# Chapter 4 Index construction

- 4.1 Hardware basics
- 4.2 Blocked sort-based indexing
- 4.3 Single-pass in-memory indexing
- 4.4 Distributed indexing
- 4.5 Dynamic indexing
- 4.6 Other types of indexes
- 4.7 References and further reading

# Outline

- 4.1 Hardware basics
- 4.2 Blocked sort-based indexing
- 4.3 Single-pass in-memory indexing
- 4.4 Distributed indexing
- 4.5 Dynamic indexing
- 4.6 Other types of indexes
- 4.7 References and further reading

## 4.1 Hardware basics

- Access to data is **much faster in memory than on disk** (roughly a factor of 10).
- **Disk seeks are “idle” time**: No data is transferred from disk while the disk head is being positioned.
- To optimize the transfer time from disk to memory: **one large chunk is faster than many small chunks**.
- **Disk I/O is block-based**: Reading and writing of entire blocks (as opposed to smaller chunks). Block sizes: 8KB to 256 KB.
- Servers used in IR systems typically have **many GBs of main memory and TBs of disk space**.
- **Fault tolerance is expensive**: It's cheaper to use many regular machines than one fault tolerant machine.

# 4.1 Hardware basics

## Some statistics (ca. 2008)

symbol	statistic	value
$s$	average seek time	5 ms = $5 \times 10^{-3}$ s
$b$	transfer time per byte	$0.02 \mu\text{s} = 2 \times 10^{-8}$ s
	processor's clock rate	$10^9 \text{ s}^{-1}$
$p$	lowlevel operation (e.g., compare & swap a word)	$0.01 \mu\text{s} = 10^{-8}$ s
	size of main memory	several GB
	size of disk space	1 TB or more

# Outline

- 4.1 Hardware basics
- 4.2 Blocked sort-based indexing
- 4.3 Single-pass in-memory indexing
- 4.4 Distributed indexing
- 4.5 Dynamic indexing
- 4.6 Other types of indexes
- 4.7 References and further reading

## 4.2 Blocked sort-based indexing

- Reuters RCV1 statistics (a dataset of text classification)
  - documents:  $N=800,000$
  - tokens per document:  $L=200$
  - terms:  $M=400,000$
  - bytes per token (incl. spaces/punctuations): 6
  - bytes per token (without spaces/punctuations): 4.5
  - bytes per term: 7.5
  - non-positional postings:  $T=100,000,000$ 
    - A posting: (termID, docID) or (termID, docID, term frequency)
- Average frequency of a term? 4.5 bytes per word token vs. 7.5 bytes per word type (i.e., term): why the difference? How many positional postings?

## 4.2 Blocked sort-based indexing

### Sort-based index construction

- As we build index, we parse documents one at a time.
- The final **postings** for any term are incomplete until the end.
- Can we keep all postings in memory and then do the sort in-memory at the end?
  - No, not for **large collections**.
  - Hence, we need to **store the intermediate results on disk**.



## 4.2 Blocked sort-based indexing

### Same algorithm for disk?

- Can we use the same index construction algorithm for larger collections, but by using disk instead of memory?
  - No, sorting very large sets of records on disk is **too slow** - too many disk seeks.
  - We need an **external** sorting algorithm (i.e., sort in the main memory).

## 4.2 Blocked sort-based indexing

### “External” sorting algorithm (using few disk seeks)

- For the RCV1 data, we must sort  $T = 100,000,000$  non-positional postings.
  - Each **posting** has size 12 bytes (4+4+4: termID, docID, term frequency) or 8 bytes (4+4: termID, docID).
- Define a **block** to consist of **10,000,000** such postings
  - We can easily fit that many postings into the main memory.
  - We will have **10** such blocks for the RCV1 data.

## 4.2 Blocked sort-based indexing

- **Basic idea**
  - For each block:
    - (i) accumulate **postings**
    - (ii) sort in the main memory w.r.t. termID
    - (iii) write to disk
  - Then **merge** the blocks into one.

## 4.2 Blocked sort-based indexing

### Algorithm: Blocked Sort-Based Indexing

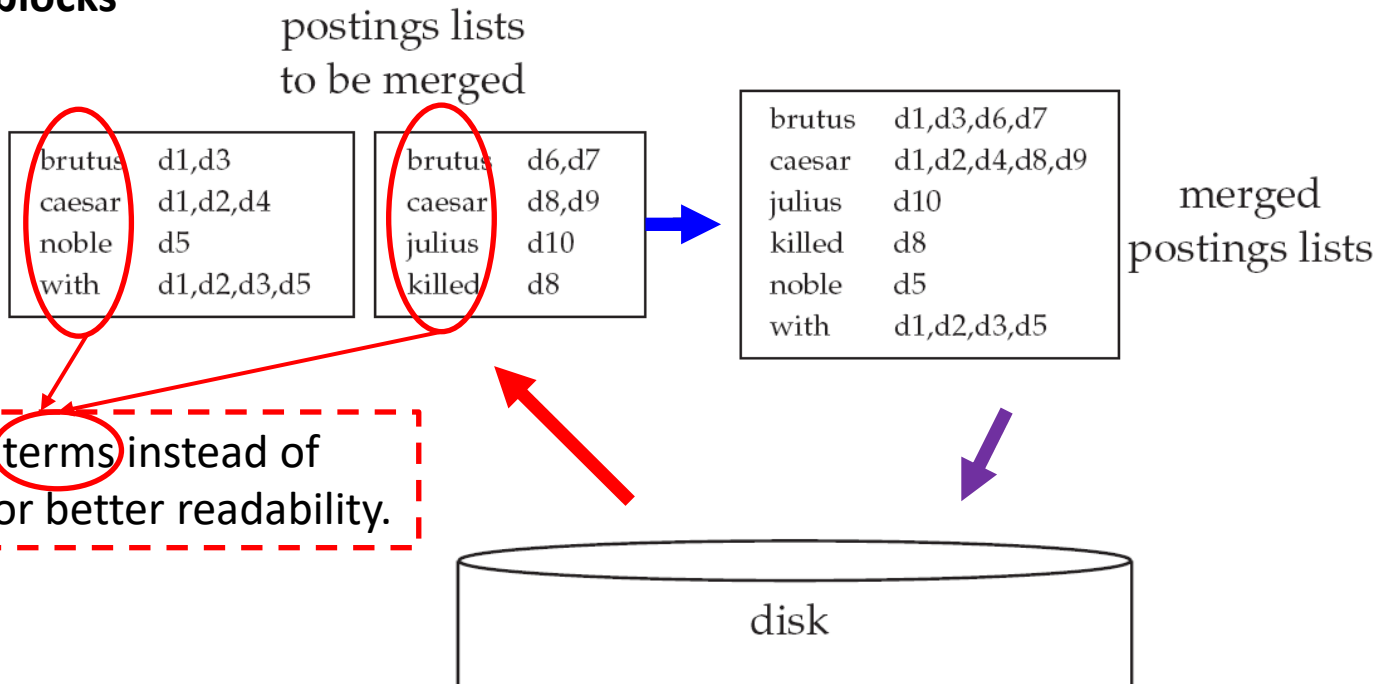
```
BSBINDEXCONSTRUCTION()  
1   $n \leftarrow 0$   
2  while (all documents have not been processed)  
3  do  $n \leftarrow n + 1$   
4      $block \leftarrow \text{PARSENEXTBLOCK}()$   
5      $\text{BSBI-INVERT}(block)$   
6      $\text{WRITEBLOCKTODISK}(block, f_n)$   
7   $\text{MERGEBLOCKS}(f_1, \dots, f_n; f_{\text{merged}})$ 
```

$n$ : the  $n$ th block

- **BSBI-Invert(block)**
  - **Sort** the (termID, docID) pairs w.r.t. termID
  - Collect all (termID, docID) pairs with the same termID into a postings list

## 4.2 Blocked sort-based indexing

### Merge blocks



Merging in blocked sort-based indexing (BSBI). Two blocks (“postings lists to be merged”) are **loaded from disk into memory**, **merged in memory** (“merged postings lists”) and **written back to disk**.

# Outline

- 4.1 Hardware basics
- 4.2 Blocked sort-based indexing
- 4.3 Single-pass in-memory indexing
- 4.4 Distributed indexing
- 4.5 Dynamic indexing
- 4.6 Other types of indexes
- 4.7 References and further reading

## 4.3 Single-pass in-memory indexing

### Problem with sort-based algorithm

- Our assumption: **We can keep the dictionary in memory**, because we need the dictionary (which grows dynamically) in order to implement **a term to termID mapping**.
- Actually, we could work with (**term**, **docID**) postings instead of (**termID**, **docID**) postings... but then intermediate files become very large (we would end up with a scalable, but very slow index construction method).

## 4.3 Single-pass in-memory indexing

### Single-pass in-memory indexing (SPIMI)

- Key idea 1(第一个想法): Generate **separate dictionaries** for each block. Then, we do not need to maintain **term-termID** mapping across blocks. Notes: each separate dictionary is **NOT** used to do term-termID mapping.
- Key idea 2(第二个想法): **Don't sort**. Accumulate postings in postings lists as they occur.
- With these two ideas, we can generate **a complete inverted index for each block**.
- These separate indexes can then be **merged** into one big index.



## 4.3 Single-pass in-memory indexing

### Algorithm

```
SPIMI-INVERT(token_stream)
1  output_file ← NEWFILE()
2  dictionary ← NEWHASH()
3  while (free memory available)
4  do token ← next(token_stream)
5      if term(token) ∉ dictionary
6          then postings_list ← ADDTODICTIONARY(dictionary, term(token))
7          else postings_list ← GETPOSTINGSLIST(dictionary, term(token))
8          if full(postings_list)
9              then postings_list ← DOUBLEPOSTINGSLIST(dictionary, term(token))
10         ADDTOPOSTINGSLIST(postings_list, docID(token))
11     sorted_terms ← SORTTERMS(dictionary)
12     WRITEBLOCKTODISK(sorted_terms, dictionary, output_file)
13     return output_file
```

token\_stream refers to the (term, docID) pairs

One separate dictionary for each block

*sorted\_terms* instead of *sorted\_termIDs*, because there is a separate dictionary for each block.

Notes: Merging of blocks is very similar to that of BSBI.

## 4.3 Single-pass in-memory indexing

### Exercise

- What is the difference between the **sort** in BSBI and the **sort** SPIMI?

## 4.3 Single-pass in-memory indexing

### **SPIMI: Compression**

- Compression makes SPIMI even more efficient.
  - Compression of terms
  - Compression of postings
  - See Chapter 5

# Outline

- 4.1 Hardware basics
- 4.2 Blocked sort-based indexing
- 4.3 Single-pass in-memory indexing
- 4.4 Distributed indexing
- 4.5 Dynamic indexing
- 4.6 Other types of indexes
- 4.7 References and further reading

## 4.4 Distributed indexing

- For web-scale indexing: must use a distributed computer cluster
- Individual machines are fault-prone. Can unpredictably slow down or fail.
- How do we exploit such a pool of machines?
  - Maintain a **master** machine directing the indexing job - considered "safe"
  - Break up indexing into sets of **parallel tasks**
  - The master machine assigns each task to an idle machine from a pool

## 4.4 Distributed indexing

### Parallel tasks

- We will define two sets of parallel tasks and deploy two types of machines to solve them:
  - **Parsers**
  - **Inverters**
- Break the input document collection into **splits** (corresponding to blocks in BSBI/SPIMI)
- Each split is a subset of documents

## 4.4 Distributed indexing

### Parsers

- **Master** assigns a split to an idle parser machine.
- **Parser** reads a document at a time and **emits** (term, docID) pairs.
- **Parser** writes pairs into **j** term-partitions.
- Each for a range of terms' first letters, e.g., a-f, g-p, q-z (here: **j = 3**)

## 4.4 Distributed indexing

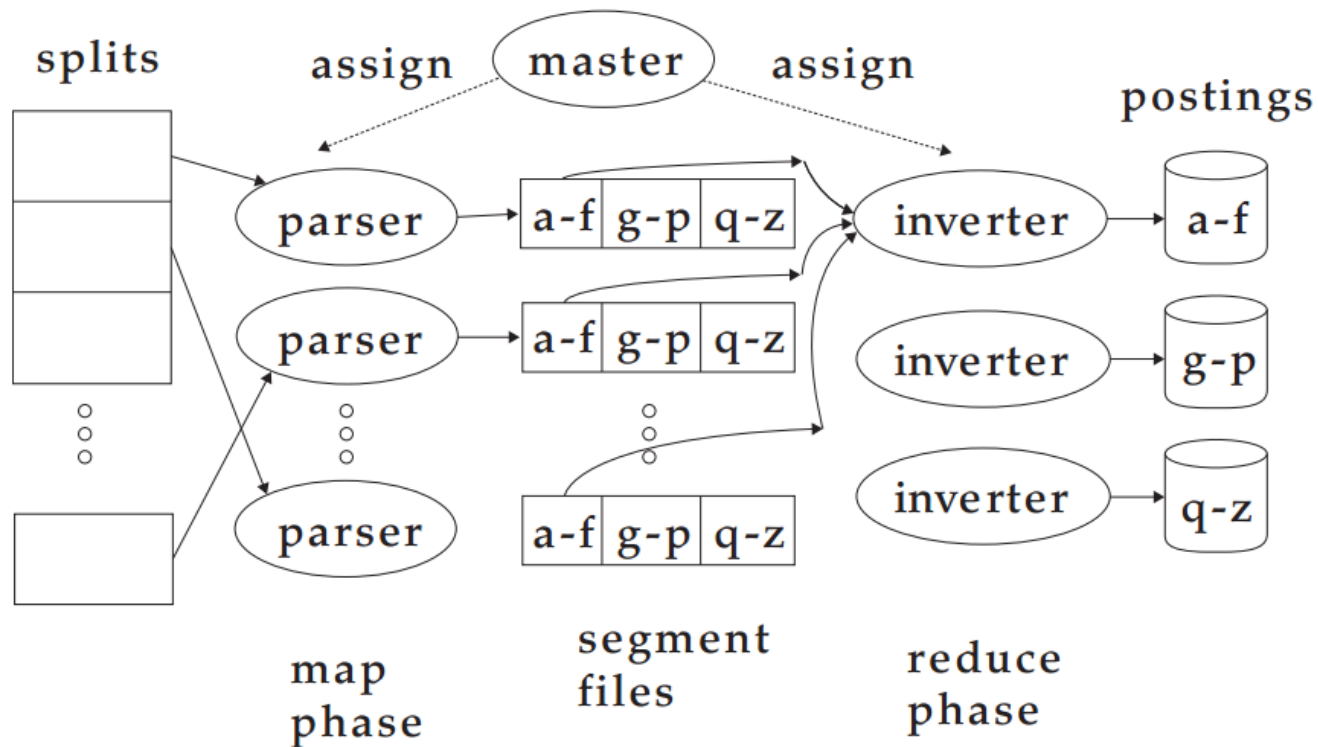
### Inverters

- An inverter collects all (term, docID) pairs (i.e., postings) for one term-partition (e.g., for a-f).
- Sorts and writes to postings lists



## 4.4 Distributed indexing

### Data flow



## 4.4 Distributed indexing

### Map/Reduce

- The index construction algorithm we just described is **an instance of Map/Reduce**.
- Map/Reduce is **a robust and conceptually simple framework** for distributed computing... without having to write code for the distribution part.
- The **Google indexing system** consisted of a number of phases, each implemented in Map/Reduce.
  - Index construction was just one phase.
  - Another phase: transform term-partitioned into document-partitioned index.

# 4.4 Distributed indexing

## Index construction in Map/Reduce

### Schema of map and reduce functions

map: input

→ list( $k, v$ )

reduce: ( $k, \text{list}(v)$ )

→ output

### Instantiation of the schema for index construction

map: web collection

→ list( $\text{termID}, \text{docID}$ )

reduce: ( $\langle \text{termID}_1, \text{list}(\text{docID}) \rangle, \langle \text{termID}_2, \text{list}(\text{docID}) \rangle, \dots$ )

→ ( $\text{postings\_list}_1, \text{postings\_list}_2, \dots$ )

### Example for index construction

map:  $d_2 : C \text{ DIED. } d_1 : C \text{ CAME, } C \text{ C'ED.}$

→ ( $\langle C, d_2 \rangle, \langle \text{DIED}, d_2 \rangle, \langle C, d_1 \rangle, \langle \text{CAME}, d_1 \rangle, \langle C, d_1 \rangle, \langle \text{C'ED}, d_1 \rangle$ )

reduce: ( $\langle C, (d_2, d_1, d_1) \rangle, \langle \text{DIED}, (d_2) \rangle, \langle \text{CAME}, (d_1) \rangle, \langle \text{C'ED}, (d_1) \rangle$ )

→ ( $\langle C, (d_1:2, d_2:1) \rangle, \langle \text{DIED}, (d_2:1) \rangle, \langle \text{CAME}, (d_1:1) \rangle, \langle \text{C'ED}, (d_1:1) \rangle$ )

# Outline

- 4.1 Hardware basics
- 4.2 Blocked sort-based indexing
- 4.3 Single-pass in-memory indexing
- 4.4 Distributed indexing
- 4.5 Dynamic indexing
- 4.6 Other types of indexes
- 4.7 References and further reading

## 4.5 Dynamic indexing

### Dynamic indexing: Simplest approach

- Maintain a big main index on disk
- New docs go into a small auxiliary index in the main memory
- Search across both, merge results
- Periodically, merge auxiliary index into big index
- Deletions:
  - Invalidation bit-vector for deleted docs
  - Filter docs returned by index using this bit-vector

## 4.5 Dynamic indexing

### Issue with auxiliary and main index

- Frequent merges
- Poor search performance during index merge

## 4.5 Dynamic indexing

### Logarithmic merge (1/3)

- Logarithmic merging amortizes the cost of merging indexes over time -> Users see smaller effect on response times.
- Maintain **a series of indexes**, each twice as large as the previous one.
  - Keep smallest (Z0) in memory
  - Larger ones (I0, I1, ...) on disk
  - If Z0 gets too big ( $>n$ ), write to disk as I0 ... or merge with I0 (if I0 already exists) and write merger to I1, etc.

## 4.5 Dynamic indexing

### Logarithmic merge (2/3)

LMERGEADDTOKEN(*indexes*,  $Z_0$ , *token*)

1  $Z_0 \leftarrow \text{MERGE}(Z_0, \{\text{token}\})$

2 **if**  $|Z_0| = n$

←  $n$  is the size of the auxiliary index

3     **then for**  $i \leftarrow 0$  **to**  $\infty$

4         **do if**  $l_i \in \text{indexes}$

5             **then**  $Z_{i+1} \leftarrow \text{MERGE}(l_i, Z_i)$

6                 ( $Z_{i+1}$  is a *temporary* index on disk.)

7                  $\text{indexes} \leftarrow \text{indexes} - \{l_i\}$

8             **else**  $l_i \leftarrow Z_i$  ( $Z_i$  becomes the *permanent* index  $l_i$ .)

9                  $\text{indexes} \leftarrow \text{indexes} \cup \{l_i\}$

10                 BREAK

11      $Z_0 \leftarrow \emptyset$

LOGARITHMICMERGE()

1  $Z_0 \leftarrow \emptyset$  ( $Z_0$  is the in-memory index.)

2  $\text{indexes} \leftarrow \emptyset$

3 **while** true

4 **do** LMERGEADDTOKEN(*indexes*,  $Z_0$ , GETNEXTTOKEN())



## 4.5 Dynamic indexing

### Logarithmic merge (3/3)

- Number of indexes bounded by  $O(\log T/n)$ , where  $n$  is the size of the auxiliary index and  $T$  is the total number of postings.
- So, query processing requires the merging of  $O(\log T/n)$  indexes.
- Time complexity of index construction is  $O(T \log T/n)$  because each of  $T$  postings is merged  $O(\log T/n)$  times.

# Summary

- 4.1 Hardware basics
- 4.2 Blocked sort-based indexing
- 4.3 Single-pass in-memory indexing
- 4.4 Distributed indexing
- 4.5 Dynamic indexing
- 4.6 Other types of indexes
- 4.7 References and further reading