

Information Retrieval

Weike Pan

The slides are **adapted from those provided by Prof. Hinrich Schütze** at University of Munich (<http://www.cis.lmu.de/~hs/teach/14s/ir/>).

Chapter 5 Index compression

- 5.1 Statistical properties of terms in information retrieval
- 5.2 Dictionary compression
- 5.3 Postings file compression
- 5.4 References and further reading

Outline

- 5.1 Statistical properties of terms in information retrieval
- 5.2 Dictionary compression
- 5.3 Postings file compression
- 5.4 References and further reading

5.1 Statistical properties of terms in information retrieval

- Reuters RCV1 statistics
 - documents: $N=800,000$
 - tokens per document: $L=200$
 - terms: $M=400,000$
 - bytes per token (including spaces/punctuations): 6
 - bytes per token (without spaces/ punctuations): 4.5
 - bytes per term: 7.5
 - non-positional postings: $T=100,000,000$

5.1 Statistical properties of terms in information retrieval

Heaps' law

- How many distinct words are there? Can we assume there is an upper bound? The vocabulary will keep growing with the collection size.
- **Heaps' law:** $M = kT^b$
 - M is the vocabulary size, T is the number of tokens in the collection.
 - Typical values for the parameters k and b are: $30 \leq k \leq 100$ and $b \approx 0.5$.
 - It is the simplest possible relationship between the collection size T and the vocabulary size M in the log-log space. It is an empirical law.
- **Question:** Can you verify the Heaps' law using the RCV1 data? (Plot the statistics in a $\log_{10} M - \log_{10} T$ space)

5.1 Statistical properties of terms in information retrieval

Zipf's law

- How many frequent vs. infrequent terms we should **expect** in a collection?
- **Zipf's law:** cf_i is the collection frequency of the i -th most frequent term t_i , i.e., the number of occurrences of the term t_i in the collection.
 - cf_i is proportional to $1/i$
- Question: Can you verify Zipf's law using the RCV1 data?

Outline

- 5.1 Statistical properties of terms in information retrieval
- 5.2 Dictionary compression
- 5.3 Postings file compression
- 5.4 References and further reading

5.2 Dictionary compression

- The dictionary is small compared with the postings file (i.e., postings lists).
- But we want to **keep it in memory**.
- Also: competition with other applications, cell phones, onboard computers, fast startup time. Hence, **compressing the dictionary is important**.

5.2 Dictionary compression

Dictionary as an array of fixed-width entries (1/2)

| term | document frequency | pointer to postings list |
|--------|--------------------|--------------------------|
| a | 656,265 | → |
| aachen | 65 | → |
| ... | ... | ... |
| zulu | 221 | → |

space needed: 20 bytes 4 bytes 4 bytes

- Space for RCV1 data: $(20+4+4)*400,000 = 11.2 \text{ MB}$


5.2 Dictionary compression

Dictionary as an array of fixed-width entries (2/2)

- Most of the bytes in the term column are **wasted**. We allot (分配) 20 bytes for terms of length 1.
- But, we **still can not handle** hydrochlorofluorocarbons and supercalifragilisticexpialidocious.
- Average length of a term in English: 8 characters (or a little bit less)
- How can we use "on average 8 characters per term"?

5.2 Dictionary compression

Dictionary as a **string** (1/2)



5.2 Dictionary compression

Dictionary as a **string** (2/2)

- 4 bytes per term for **document frequency**
- 4 bytes per term for **pointer to postings list**
- 8 bytes (on average) for **term** in string
- **3 bytes** per pointer into string (needs $\log_2(8 * 400000) = 21.6 < 24$ **bits** to resolve $8 * 400000$ positions)
- Space: $400000 * (4 + 4 + 8 + 3) = 7.6\text{MB}$ (compared to 11.2 MB for fixed-width array)

5.2 Dictionary compression

Dictionary as a string with **blocking** (1/2)

...7systile9syzygetic8syzygial6syzygy11szaibelyite6szecin...

| freq. | postings ptr. | term ptr. |
|-------|---------------|-----------|
|-------|---------------|-----------|

9

→

92

→

5

→

71

→


12

→

...

...

...



5.2 Dictionary compression

Dictionary as a string with **blocking** (2/2)

- Example: block size $k = 4$
 - we use 4×3 bytes for term pointers without blocking ...
 - we now use **3 bytes** for one pointer plus **4 bytes** for indicating the length of each term.
 - We **save $12 - (3 + 4) = 5$ bytes per block.**
 - Total savings: $400000/4 * 5 = \mathbf{0.5\ MB}$
- This reduces the size of the dictionary from 7.6 MB to **7.1 MB**.

5.2 Dictionary compression

Front coding

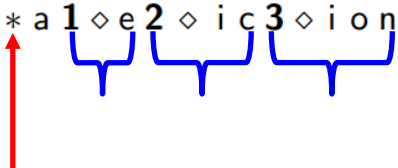
One block in blocked compression ($k = 4$) ...

8 a u t o m a t a **8** a u t o m a t e **9** a u t o m a t i c **10** a u t o m a t i o n



... further compressed with front coding.

8 a u t o m a t * a **1** ◇ e **2** ◇ i c **3** ◇ i o n

A diagram illustrating front coding. A red arrow points to the asterisk in 'mat*a'. Blue brackets are placed under the characters 'e', 'i', and 'i' in the subsequent words, indicating they are stored as offsets from the dictionary entry 'mat*a'.

5.2 Dictionary compression

Dictionary compression for the RCV1 data: Summary

| data structure | size in MB |
|---------------------------------------|------------|
| dictionary, fixed-width | 11.2 |
| dictionary, term pointers into string | 7.6 |
| ~, with blocking, $k = 4$ | 7.1 |
| ~, with blocking & front coding | 5.9 |

Outline

- 5.1 Statistical properties of terms in information retrieval
- 5.2 Dictionary compression
- 5.3 Postings file compression
- 5.4 References and further reading

5.3 Postings file compression

- The postings file (i.e., postings list) is much larger than the dictionary: factor of at least 10.
- A posting for our purpose is a **docID**.
- For the RCV1 data (800000 documents), we would use **32 bits** per docID when using 4-byte integers.
- Alternatively, we can use $\log_2 \underline{800000} \approx 19.6 < \mathbf{20 \text{ bits}}$ per docID.
- Our goal: use **a lot less than 20 bits** per docID.

5.3 Postings file compression

Key idea: Store **gaps instead of docIDs**

- Each postings list is ordered **in increasing order** of docID.
 - Example postings list: computer: 283154, 283159, 283202, ...
- It suffices to store **gaps**: $283159 - 283154 = 5$, $283202 - 283159 = 43$
- Example postings list using gaps: computer: 283154, **5**, **43**, ...
- Gaps for frequent terms are **small**.
- Thus: We can **encode small gaps** with **fewer than 20 bits**.

5.3 Postings file compression

Gap encoding

| | encoding | postings list | | | | | |
|----------------|----------|---------------|--------|--------|--------|--------|-----|
| THE | docIDs | ... | 283042 | 283043 | 283044 | 283045 | ... |
| | gaps | | 1 | 1 | 1 | | ... |
| COMPUTER | docIDs | ... | 283047 | 283154 | 283159 | 283202 | ... |
| | gaps | | 107 | 5 | 43 | | ... |
| ARACHNOCENTRIC | docIDs | 252000 | 500100 | | | | |
| | gaps | 252000 | 248100 | | | | |

5.3 Postings file compression

Variable length encoding

- Aim:
 - For ARACHNOCENTRIC and other **rare** terms, we will use **about 20 bits per gap**.
 - For THE and other very **frequent** terms, we will use **only a few bits per gap**.
- In order to implement this, we need to devise some form of variable length encoding
 - **Variable length encoding** uses few bits for small gaps and many bits for large gaps.

5.3 Postings file compression

Variable byte (VB) code

- Used by many commercial/research systems
- Dedicate one bit (high bit) to be **a continuation bit** c .
 - **If** the gap G fits within 7 bits, binary-encode it in the 7 available bits and set **$c = 1$** .
 - **Else**: encode lower-order 7 bits and then use one or more additional bytes to encode the higher order bits using the same algorithm.
- At the end, set the **continuation bit** of **the last byte to 1 ($c = 1$)** and of **the other bytes to 0 ($c = 0$)**.

5.3 Postings file compression

VB code encoding algorithm

```
VBENCODENUMBER(n)  
1  bytes  $\leftarrow \langle \rangle$   
2  while true  
3  do PREPEND(bytes, n mod 128)  
4    if n < 128  
5      then BREAK  
6    n  $\leftarrow n \text{ div } 128$   
7  bytes[LENGTH(bytes)] += 128  
8  return bytes
```

```
VBENCODE(numbers)  
1  bytestream  $\leftarrow \langle \rangle$   
2  for each n  $\in$  numbers  
3  do bytes  $\leftarrow$  VBENCODENUMBER(n)  
4    bytestream  $\leftarrow$  EXTEND(bytestream, bytes)  
5  return bytestream
```

- PREPEND: adds an element to the beginning of a list, e.g., PREPEND(< 1,2 >,3) = < 3,1,2 >.
- *bytes*[LENGTH(*bytes*)] += 128: continuation bit (c=1)
- EXTEND: extends a list, e.g., EXTEND(<1,2>, <3,4>) = <1,2,3,4>.

5.3 Postings file compression

VB code decoding algorithm

```
VBDECODE(bytestream)
1  numbers  $\leftarrow \langle \rangle$ 
2  n  $\leftarrow 0$ 
3  for i  $\leftarrow 1$  to LENGTH(bytestream)
4  do if bytestream[i] < 128
5      then n  $\leftarrow 128 \times n + \textit{bytestream}[\textit{i}]$ 
6      else n  $\leftarrow 128 \times n + (\textit{bytestream}[\textit{i}] - 128)$ 
7          APPEND(numbers, n)
8          n  $\leftarrow 0$ 
9  return numbers
```


Diagram illustrating the VB code decoding algorithm:

- Line 4: *bytestream*[*i*] < 128 is checked. A purple arrow points from this condition to a green box labeled "continuation bit (c=0)".
- Line 6: (*bytestream*[*i*] - 128) is calculated. A red arrow points from this expression to a green box labeled "continuation bit (c=1)".

5.3 Postings file compression

Example

| | | | |
|---------|--------------------------|-----------------|-----------------------------------|
| docIDs | 824 | 829 | 215406 |
| gaps | | 5 | 214577 |
| VB code | <u>00000110 10111000</u> | <u>10000101</u> | <u>00001101 00001100 10110001</u> |



5.3 Postings file compression

Other variable codes

- Instead of bytes, we can also use a different “unit of alignment”: 32 bits (words), 16 bits, 4 bits (nibbles), etc
- Variable byte alignment wastes space if you have many small gaps - nibbles do better on those.

5.3 Postings file compression

Gamma codes for gap encoding

- [illegible]

5.3 Postings file compression

Gamma code

- Represent a gap G as a pair of **length** and **offset**.
 - Offset is the gap in binary, **with the leading bit chopped off**, e.g., 13→1101→**101** = offset.
 - Length is the length of offset, e.g, for 13 (the offset is 101), this is 3. Encode length in **unary** code: **1110**.
 - Gamma code of 13 is the concatenation of length and offset: **1110****101**.

5.3 Postings file compression

Gamma code examples

| number | unary code | length | offset | γ code |
|--------|------------|-------------|------------|------------------------|
| 0 | 0 | | | |
| 1 | 10 | 0 | | 0 |
| 2 | 110 | 10 | 0 | 10,0 |
| 3 | 1110 | 10 | 1 | 10,1 |
| 4 | 11110 | 110 | 00 | 110,00 |
| 9 | 1111111110 | 1110 | 001 | 1110,001 |
| 13 | | 1110 | 101 | 1110,101 |
| 24 | | 11110 | 1000 | 11110,1000 |
| 511 | | 1111111110 | 11111111 | 111111110,11111111 |
| 1025 | | 11111111110 | 0000000001 | 11111111110,0000000001 |

5.3 Postings file compression

Length of gamma code

- The length of **offset** is $\lceil \log_2 G \rceil$ bits.
- The length of **length** is $\lceil \log_2 G \rceil + 1$ bits,
- So the length of the entire code is $2 * \lceil \log_2 G \rceil + 1$ bits.
- Gamma codes are always of odd length.
- Gamma codes are within a factor of 2 of the optimal encoding length $\log_2 G$.

5.3 Postings file compression

Gamma code: Properties

- Gamma code (like variable byte code) is **prefix-free**: a valid code word is not a prefix of any other valid code.
- Gamma code is **parameter-free**.
- More theoretical analysis on the compression rate can be found in the textbook.

5.3 Postings file compression

Gamma codes: Alignment

- Machines have word boundaries, e.g., 8, 16, 32 bits.
- Compressing and manipulating at granularity of bits can be slow.
- Variable byte encoding is aligned and thus potentially more efficient.
- Regardless of efficiency, variable byte is conceptually simpler at little additional space cost.

5.3 Postings file compression

Compression of the RCV1 data

| data structure | size in MB |
|---------------------------------------|------------|
| dictionary, fixed-width | 11.2 |
| dictionary, term pointers into string | 7.6 |
| ~, with blocking, $k = 4$ | 7.1 |
| ~, with blocking & front coding | 5.9 |
| collection (text, xml markup etc) | 3600.0 |
| collection (text) | 960.0 |
| T/D incidence matrix | 40,000.0 |
| postings, uncompressed (32-bit words) | 400.0 |
| postings, uncompressed (20 bits) | 250.0 |
| postings, variable byte encoded | 116.0 |
| postings, γ encoded | 101.0 |

5.3 Postings file compression

Summary

- We can now create an index for highly efficient Boolean retrieval that is **very space efficient**.
- Only 10-15% of the total size of the text in the collection.
- However, we've ignored **positional** and **frequency** information.
- For this reason, space savings are less in reality.

Summary

- 5.1 Statistical properties of terms in information retrieval
- 5.2 Dictionary compression
- 5.3 Postings file compression
- 5.4 References and further reading