

Information Retrieval

Weike Pan

The slides are **adapted from those provided by Prof. Hinrich Schütze** at University of Munich (<http://www.cis.lmu.de/~hs/teach/14s/ir/>).

Chapter 6 Scoring, term weighting & the vector space model

- 6.1 Parametric and zone indexes
- 6.2 Term frequency and weighting
- 6.3 The vector space model for scoring
- 6.4 Variant tf-idf functions
- 6.5 References and further reading

Outline

- 6.1 Parametric and zone indexes
- 6.2 Term frequency and weighting
- 6.3 The vector space model for scoring
- 6.4 Variant tf-idf functions
- 6.5 References and further reading

6.1 Parametric and zone indexes

- Thus far, our queries have been **Boolean**. Documents either match or don't.
- Good for **expert users** with precise understanding of their needs and of the collection.
- Good for **applications**: Applications can easily consume 1000s of results.
- **Not good for the majority of users**. Most users **are not capable of writing Boolean queries** ... or they are, but they think it's too much work. Most users don't want to wade through 1000s of results. This is particularly true of web search.

6.1 Parametric and zone indexes

Problem with Boolean search: Feast or famine

- Boolean queries often result in either **too few** (=0) or **too many** (1000s) results.
- In Boolean retrieval, it takes a lot of skill to come up with a query that produces **a manageable number of hits**.

6.1 Parametric and zone indexes

Feast or famine: No problem in ranked retrieval

- With **ranking**, large result sets are not an issue.
- Just show the **top 10** results.
- Doesn't overwhelm the user.
- Premise (前提): the ranking algorithm works, i.e., **more relevant results are ranked higher than less relevant results**

6.1 Parametric and zone indexes

Scoring as the basis of ranked retrieval

- How can we accomplish a **relevance ranking** of the documents with respect to a query?
- Assign a **score** to each **query-document pair**, say in $[0,1]$. This **score** measures how well the document and the query "**match**".
- **Sort** the documents according to the **scores**.

6.1 Parametric and zone indexes

Query-document matching scores

- How do we compute the score of a query-document pair?
 - If **no query term occurs** in the document: score should be 0.
 - The **more frequent a query term** in the document, the higher the score. (Notes: term frequency)
 - The **more query terms** occur in the document, the higher the score. (Notes: number of terms)

6.1 Parametric and zone indexes

Bag of words model

- We do not consider the **order** of words in a document.
- *John is quicker than Mary* and *Mary is quicker than John* are represented exactly the same way.
- In a sense, this is a step back: The positional index (see chapter 2) is able to distinguish these two documents.

Outline

- 6.1 Parametric and zone indexes
- 6.2 Term frequency and weighting
- 6.3 The vector space model for scoring
- 6.4 Variant tf-idf functions
- 6.5 References and further reading

6.2 Term frequency and weighting

TF (Term Frequency)

- Term frequency: number of occurrences of a term t in a document d
- Raw term frequency is not what we want because: A document with $tf = 10$ occurrences of the term is more relevant than a document with $tf = 1$ occurrence of the term. But not 10 times more relevant, i.e., the relevance does not increase proportionally with the term frequency.
- We define the **TF weight** of term t in document d as follows

$$w_{t,d} = \begin{cases} 1 + \log_{10} tf_{t,d} & \text{if } tf_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases}$$

6.2 Term frequency and weighting

IDF (Inverse Document Frequency)

- Rare terms in the collection (instead of in a document as that in TF) are more informative than frequent terms.
 - high weights for rare terms
 - low (positive) weights for frequent terms
- We define the **IDF weight** of term **t** in the **collection** as follows

$$\text{idf}_t = \log_{10} \frac{N}{\text{df}_t}$$

- The **IDF weight** is a measure of the informativeness of the term in the collection. Note that the **document frequency** denotes the number of documents in the collection that the term occurs in.

6.2 Term frequency and weighting

TF-IDF

- The **TF-IDF weight** of a term **t** in a document **d** is the product of the corresponding **TF** weight and **IDF** weight:

$$w_{t,d} = (1 + \log \text{tf}_{t,d}) \cdot \log \frac{N}{\text{df}_t}$$

- Increases with the number of **occurrences** within a **document** (TF).
- Increases with the **rarity** of the term in the **collection** (IDF).
- The TF-IDF weight is the best known weighting scheme in information retrieval.
- We can also see that the **IDF weight affects the ranking of documents for queries with at least two terms**. It has little effect on document ranking for one-term queries.

Outline

- 6.1 Parametric and zone indexes
- 6.2 Term frequency and weighting
- 6.3 The vector space model for scoring
- 6.4 Variant tf-idf functions
- 6.5 References and further reading

6.3 The vector space model for scoring

Ranked retrieval in the vector space model

- Step 1. Represent the query as a weighted TF-IDF **vector**
- Step 2. Represent each document as a weighted TF-IDF **vector**
- Step 3. Compute the **cosine similarity** between the query **vector** and each document **vector**

$$\cos(\vec{q}, \vec{d}) = \text{SIM}(\vec{q}, \vec{d}) = \frac{\vec{q} \cdot \vec{d}}{|\vec{q}| |\vec{d}|} = \frac{\sum_{i=1}^{|V|} q_i d_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2} \sqrt{\sum_{i=1}^{|V|} d_i^2}}$$

- Step 4. Rank the documents with respect to the query
- Step 5. Return the top K (e.g., K = 10) documents to the user

Summary

- 6.1 Parametric and zone indexes
- 6.2 Term frequency and weighting
- 6.3 The vector space model for scoring
- 6.4 Variant tf-idf functions
- 6.5 References and further reading